

# Database Conditionals, Grouping, and Joins



**MIT** Center for  
Transportation & Logistics

# Motivating questions

- How can we narrow the sets of records that get returned from a query?
- How can we make statistical queries across different groupings of records?
- How can we sample or order our output results?
- How can we integrate data from another source with our database?

# Narrowing select statements with conditional clauses

# WHERE IN

- WHERE *attribute* IN is used to select rows that are satisfied by a set of WHERE conditions on the same attribute:

```
SELECT *  
FROM Offices  
WHERE State IN ('CO', 'UT', 'TX');
```

- Equivalent to:

```
SELECT *  
FROM Offices  
WHERE State = 'CO' OR State = 'UT'  
      OR State = 'TX';
```

OfficeNbr	City	State	Region	Target	Sales	Phone
1	Denver	CO	West	3000000.00	130000.00	970.586.3341
57	Dallas	TX	West	0.00	0.00	214.781.5342

# BETWEEN keyword

- Select records where the attribute value is between two numbers using BETWEEN
- Range is inclusive and also works with time and date data

```
SELECT *  
FROM SalesReps  
WHERE Quota BETWEEN 50000 AND 100000;
```

RepNbr	Name	RepOffice	Quota	Sales
53	Bill Smith	1	100000.00	0.00
89	Jen Jones	2	50000.00	130000.00

```
SELECT *  
FROM SalesReps  
WHERE Quota NOT BETWEEN 50000 AND 100000;
```

RepNbr	Name	RepOffice	Quota	Sales
--------	------	-----------	-------	-------

# NULL values

- Empty or missing values are stored in tables as `NULL`
- `NULL` values evaluate to `NOT TRUE` in all cases
- Insert a new Sales Representative with a `NULL` Quota:

```
INSERT INTO SalesReps  
VALUES ('25', 'Chris', '1', NULL, 1000.0)
```

RepNbr	Name	RepOffice	Quota	Sales
53	Bill Smith	1	100000.00	0.00
89	Jen Jones	2	50000.00	130000.00
25	Chris	1	NULL	1000.00

# NULLs

- The following two sets of queries will not return all sales reps:

```
SELECT Name
FROM SalesReps
WHERE Sales > Quota;
```

```
SELECT Name
FROM SalesReps
WHERE Sales = Quota;
```

```
SELECT Name
FROM SalesReps
WHERE Sales <= Quota;
```

```
SELECT Name
FROM SalesReps
WHERE Sales <> Quota;
```

RepNbr	Name	RepOffice	Quota	Sales
53	Bill Smith	1	100000.00	0.00
89	Jen Jones	2	50000.00	130000.00
25	Chris	1	NULL	1000.00

# NULLs

- Check for NULLS using IS:

```
SELECT Name  
FROM SalesReps  
WHERE Quota IS NULL;
```

- Check for NULLS using IS NOT:

```
SELECT Name  
FROM SalesReps  
WHERE Quota IS NOT NULL;
```



# Key points from lesson

- **WHERE IN** statements are used to identify records in a table with an attribute **matching a value from a specified set of values**
- **BETWEEN** keywords are used to identify records that have **values for a particular attribute that fall within a specified range**
- When a specific attribute may contain **NULL or missing values**, special care must be taken when using these **conditional clauses**

# Grouping Data and Statistical Functions

# GROUP BY with COUNT(\*)

- Find the number of sales for each customer:

```
SELECT Cust, COUNT(*)  
FROM Orders  
GROUP BY Cust;
```

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

Cust	COUNT(*)
211	1
522	2

# GROUP BY with COUNT(\*) and HAVING

- Return the number of parts from each vendor:

```
SELECT Vendor, COUNT(*)  
FROM Parts  
GROUP BY Vendor;
```

Vendor	COUNT(*)
A	4
B	1
C	2

- Return the number of parts from each vendor who sells more than two parts:

```
SELECT Vendor, COUNT(*)  
FROM Parts  
GROUP BY Vendor  
HAVING COUNT(*) > 2;
```

Vendor	COUNT(*)
A	4

PartID	Vendor
123	A
234	A
345	B
362	A
2345	C
3464	A
4533	C

# Aggregate Statistical Functions in SQL

- Statistical functions are available in SQL to perform analytics
- Commonly used functions include:

Name	Description
AVG()	Return the average value of the argument
COUNT()	Return a count of the number of rows returned
COUNT(DISTINCT)	Return the count of a number of different values
MAX()	Return the maximum value
MIN()	Return the minimum value
STD()	Return the population standard deviation
STDDEV_SAMP()	Return the sample standard deviation
SUM()	Return the sum
VAR_POP()	Return the population standard variance
VAR_SAMP()	Return the sample variance
VARIANCE()	Return the population standard variance

# Advanced Statistical Functions in SQL

- More advanced statistical functions can be created using the basic statistical functions built into SQL
- Calculate the **weighted average of student scores** from the table of **grades**:

```
SELECT SUM(student_score*score_weight) /  
SUM(score_weight)  
FROM grades;
```

- Get the z-score values for student\_score from the table grades:

```
SELECT (student_score-AVG(student_score)) /  
VARIANCE(student_score)  
FROM grades;
```

# Key points from lesson

- Built-in **statistical functions** exist in many implementations of SQL
- Statistical functions can operate on a **group of records** and return a **single value** for each group of records created by a **GROUP BY** clause
- To **restrict the output of a GROUP BY clause** to results which meet specific conditions, use the **HAVING** keyword, which is analogous to a **WHERE** clause

# Sorting and Sampling Data



# ORDER BY

- Order by one column, ascending, ASC, or descending, DESC

```
SELECT *  
FROM Customers  
ORDER BY Country DESC;
```

- Order by two columns, one first, and then the next:

```
SELECT *  
FROM Customers  
ORDER BY Country ASC, CustomerName DESC;
```

# LIMIT the number of returned records

- Example to return 5 people from the Persons table:

```
SELECT *  
FROM Persons  
ORDER BY Last_Name ASC  
LIMIT 5;
```

# Randomly select and order records

- RAND function can generate random numbers for various uses
- Reorder the entire table:

```
SELECT *  
FROM table  
ORDER BY RAND();
```

- Randomly select a single record:

```
SELECT *  
FROM table  
ORDER BY RAND()  
LIMIT 1;
```

# Randomly select and order records

- Generate a random number in the output results:

```
SELECT id, RAND()  
FROM table;
```

# Key points from lesson

- The **ORDER BY** clause specifies that the results from a query should be returned in **ascending** or **descending** order
- A **LIMIT** clause **restricts the number of records** that would be returned to a subset, which can be convenient for inspection or efficiency
- The **RAND()** function can be used to **generate random values** in the output or to **randomly sample** or **randomly order** the results of a query

# Creating New Tables and Aliases

# AS Keyword (Aliases)

- AS can be used to create an alias for a field
- A new composite address field can be created from all of the address components as:

```
SELECT CustomerName, CONCAT(Address, ', ',  
City, ', ', PostalCode, ', ', Country)  
AS Address  
FROM Customers;
```

# CREATE TABLE AS

- Use CREATE TABLE with AS to create a new table in the database using a select query

```
CREATE TABLE new_table
AS ( SELECT column_name(s)
      FROM old_table );
```

- Matches columns and data types based on the results in the select statement

```
CREATE TABLE CustomersBackup2017
AS ( SELECT CustomerName, ContactName
      FROM Customers );
```



# SELECT INTO

- Take the results of a select statement and put them in an existing table or database:

```
SELECT column_name(s)  
INTO newtable [IN externaldb]  
FROM table1;
```

- Example:

```
SELECT CustomerName, ContactName  
INTO CustomersBackup2017  
FROM Customers;
```

# Key points from lesson

- The **AS** keyword creates an alias for an attribute or result of a function that is returned in a query
- Results from a query can be inserted into a new table using the **CREATE TABLE** with the **AS** keyword
- Results from a query can be inserted into an existing table using a **SELECT INTO** clause if the table with the appropriate structure already exists

# Joining Multiple Tables

# Introduction to JOINS

- Last week, our focus was on taking large datasets and normalizing them by separating the data into different tables
- This week, the focus is on taking data from different tables and combining it together
- This may include data from other data sources which complement our own:
  - demographic information for a zip code
  - price structure for shipping zones for a carrier
- The process of merging two separate tables is called “joining”

# JOIN details

- Joins effectively merge two tables together based on a relationship between different columns in each table
- Relationships become explicit when you query the database
- Both tables will still contain the normal separated data, and data can still be added to the individual tables

# Columns in a JOIN

- Joins may be done on any columns in two tables, as long as the merge operation makes logical sense
  - They don't need to be keys, though they usually are
  - Join columns must have compatible data types
  - Join column is usually key column: Either primary or foreign
  - NULLs will never join

# JOIN Example

- List all orders, showing order number and amount along with the name and credit limit of each customer
  - Orders table has order number and amount, but no customer names or credit limits
  - Customers has customer names and credit limit, but no order info
- Cust = CustNbr

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

CustNbr	Company	CustRep	CreditLimit
211	Connor Co	89	50000.00
522	Amaratunga Enterprises	89	40000.00
890	Feni Fabricators	53	1000000.00

# JOIN Example

- List all orders, showing order number and amount along with the name and credit limit of each customer

```
SELECT OrderNbr, Amt, Company, CreditLimit
FROM Customers, Orders
WHERE Cust = CustNbr; -- (Implicit syntax)
```

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

CustNbr	Company	CustRep	CreditLimit
211	Connor Co	89	50000.00
522	Amaratunga Enterprises	89	40000.00
890	Feni Fabricators	53	1000000.00



# JOIN Example

- List all orders, showing order number and amount along with the name and credit limit of each customer

```
SELECT OrderNbr, Amt, Company, CreditLimit
FROM Customers, Orders
WHERE Cust = CustNbr; -- (Implicit syntax)
```

OrderNbr	Amt	Company	CreditLimit
1	31000.00	Connor Co	50000.00
2	4000.00	Amaratunga Enterprises	40000.00
3	500000.00	Amaratunga Enterprises	40000.00

# JOIN Example

- List all orders, showing order number and amount along with the name and credit limit of each customer

```
SELECT OrderNbr, Amt, Company, CreditLimit
FROM Customers INNER JOIN Orders
    ON Customers.CustNbr = Orders.Cust;
-- (SQL-92 syntax)
```

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

CustNbr	Company	CustRep	CreditLimit
211	Connor Co	89	50000.00
522	Amaratunga Enterprises	89	40000.00
890	Feni Fabricators	53	1000000.00


# JOIN Example

- List all orders, showing order number and amount along with the name and credit limit of each customer

```
SELECT OrderNbr, Amt, Company, CreditLimit
FROM Customers INNER JOIN Orders
    ON Customers.CustNbr = Orders.Cust;
-- (SQL-92 syntax)
```

OrderNbr	Amt	Company	CreditLimit
1	31000.00	Connor Co	50000.00
2	4000.00	Amaratunga Enterprises	40000.00
3	500000.00	Amaratunga Enterprises	40000.00

# Joins - visually



SELECT \*

FROM tb1, tb2;

Blue	Grey	Blue	Pink
Blue	Grey	Blue	Pink
Blue	Grey	Green	Red
Green	Dark Grey	Blue	Pink
Green	Dark Grey	Blue	Pink
Green	Dark Grey	Green	Red

SELECT \*

FROM tb1, tb2

WHERE tb1.bg = tb2.bg;

Blue	Grey	Blue	Pink
Blue	Grey	Blue	Pink
Green	Dark Grey	Green	Red

SELECT grey, pink

FROM tb1, tb2

WHERE tb1.bg = tb2.bg;

Grey	Pink
Grey	Pink
Dark Grey	Red

Two tables with shared field/data

SELECT \*

FROM tb1, tb2;

SELECT \*

FROM tb1, tb2

WHERE tb1.bg = tb2.bg;

SELECT grey, pink

FROM tb1, tb2

WHERE tb1.bg = tb2.bg;

# Further notes on JOINS

- Use `*` carefully in joins – it returns all columns from all tables being joined
- If a field has the same name in the tables being joined, specify the table name along with the field name:
  - `table1.fieldname, table2.fieldname`
  - `Customers.CustNbr, Orders.Amt`

# Key points from lesson

- The relational database model allows us join multiple tables to build new and unanticipated relationships
- The columns in a join must be of matching types and also must represent the same concept in two different tables
- This can help us to contextualize or integrate a table in our database with data from an external source

# Types of JOINS and VIEWS

# Join with 3 tables

```
SELECT OrderNbr, Amt, Company, Name
FROM Orders, Customers, SalesReps
WHERE Cust = CustNbr AND CustRep = RepNbr AND
      Amt > 25000;
-- (Implicit syntax)
```

- List orders over \$25,000, including the name of the salesperson who took the order and the name of the customer who placed the order

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

CustNbr	Company	CustRep	CreditLimit
211	Connor Co	89	50000.00
522	Amaratunga Enterprises	89	40000.00
890	Feni Fabricators	53	1000000.00

RepNbr	Name	RepOffice	Quota	Sales
53	Bill Smith	1	100000.00	0.00
89	Jen Jones	2	50000.00	130000.00



# Join with 3 tables

```
SELECT OrderNbr, Amt, Company, Name
FROM Orders, Customers, SalesReps
WHERE Cust = CustNbr AND CustRep = RepNbr AND
      Amt > 25000;
-- (Implicit syntax)
```

- List orders over \$25,000, including the name of the salesperson who took the order and the name of the customer who placed the order

OrderNbr	Amt	Company	Names
1	31000.00	Connor Co	Jen Jones
3	500000.00	Amaratunga Enterprises	Jen Jones

# Join with 3 tables

```
SELECT OrderNbr, Amt, Company, Name
FROM SalesReps INNER JOIN Customers
    ON SalesReps.RepNbr = Customers.CustRep INNER JOIN Orders
    ON Customers.CustNbr = Orders.Cust
WHERE Amt > 25000;
```

-- (SQL-92 syntax)

- List orders over \$25,000, including the name of the salesperson who took the order and the name of the customer who placed the order

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

CustNbr	Company	CustRep	CreditLimit
211	Connor Co	89	50000.00
522	Amaratunga Enterprises	89	40000.00
890	Feni Fabricators	53	1000000.00

RepNbr	Name	RepOffice	Quota	Sales
53	Bill Smith	1	100000.00	0.00
89	Jen Jones	2	50000.00	130000.00

# Join with 3 tables

```
SELECT OrderNbr, Amt, Company, Name
FROM SalesReps INNER JOIN Customers
      ON SalesReps.RepNbr = Customers.CustRep INNER JOIN Orders
      ON Customers.CustNbr = Orders.Cust
WHERE Amt > 25000;

-- (SQL-92 syntax)
```

- List orders over \$25,000, including the name of the salesperson who took the order and the name of the customer who placed the order

OrderNbr	Amt	Company	Names
1	31000.00	Connor Co	Jen Jones
3	500000.00	Amaratunga Enterprises	Jen Jones

# JOIN types

```
SELECT OrderNbr, Amt, Company,  
CreditLimit FROM Customers, Orders  
WHERE Cust = CustNbr;
```

- **INNER JOIN**: returns only the records with matching keys (joins common column values)
- **LEFT JOIN**: returns all rows from **LEFT** (first) table, whether or not they match a record in the second table
- **RIGHT JOIN**: returns all rows from **RIGHT** (second) table, whether or not they match a record in the first table
- **OUTER JOIN**: Returns all rows from both tables, whether or not they match (Microsoft SQL, not MySQL)
- In MySQL, **JOIN** and **INNER JOIN** are equivalent

# VIEWS

- VIEWS are virtual tables which present data in a denormalized form to users
- VIEWS DO NOT create separate copies of the data; they reference the data in the underlying tables
- Database stores a definition of a view; the data is updated each time the VIEW is invoked
- Advantages:
  - User queries are simpler on views constructed for them
  - Security: can restrict access to data in views for users
  - Independence: user or program are not affected by small changes in underlying tables

# VIEWS

```
CREATE VIEW CustomerOrders AS
SELECT CustNbr, Company, Name, OrderNbr, Prod, Qty, Amt
FROM Customers, SalesReps, Orders
WHERE CustRep = RepNbr AND CustNbr = Cust;
-- (Implicit syntax)
```

OrderNbr	Cust	Prod	Qty	Amt	Disc
1	211	Bulldozer	7	31000.00	0.20
2	522	Riveter	2	4000.00	0.30
3	522	Crane	1	500000.00	0.40

CustNbr	Company	CustRep	CreditLimit
211	Connor Co	89	50000.00
522	Amaratunga Enterprises	89	40000.00
890	Feni Fabricators	53	1000000.00

RepNbr	Name	RepOffice	Quota	Sales
53	Bill Smith	1	100000.00	0.00
89	Jen Jones	2	50000.00	130000.00

# VIEWS

```
CREATE VIEW CustomerOrders AS
SELECT CustNbr, Company, Name, OrderNbr, Prod, Qty, Amt
FROM Customers, SalesReps, Orders
WHERE CustRep = RepNbr AND CustNbr = Cust;
-- (Implicit syntax)
```

CustNbr	Company	Name	OrderNbr	Prod	Qty	Amt
211	Connor Co	Jen Jones	1	Bulldozer	7	31000.00
522	Amaratunga Enterprises	Jen Jones	2	Riveter	2	4000.00
522	Amaratunga Enterprises	Jen Jones	3	Crane	1	500000.00

# Key points from lesson

- Joining three tables together just involves one additional join between two already joined tables and a third table
- Different types of joins can be used to merge two tables together that always include every row in the left table, right table, or in both tables
- Views are virtual tables which do not change the underlying data but can be helpful to generate reports and simplify complicated queries